



Parallel Hardware

Parallel Computing – COSC 3P93

Course Outline

- Introduction
- Parallel Hardware
- Parallel Software
- Performance
- Parallel Programming Models
- Patterns
- SMP with Threads and OpenMP
- Distributed-Memory Programming with MPI
- Algorithms
- Tools
- Parallel Program Design

Today Class

- Parallel Hardware
 - Flynn's Taxonomy
 - SIMD systems
 - MIMD systems

Parallel Hardware

- Multiple issue and pipelining
 - Replication of functional units
 - Both
 - Extensions of the von Neumann model
 - Visible to programmers
 - Aware of them for better program design

Flynn's Classical Taxonomy

- There are different ways to classify parallel computers
- Flynn's Taxonomy → 1966
 - One of the more widely used classifications
- It distinguishes multi-processor computer architectures according to how they can be classified along the two independent dimensions of Instruction Stream and Data Stream
 - Each of these dimensions can have only one of two possible states
 - Single
 - Or multiple
- Classification of computer architectures
 - Simultaneous coordination
 - # of instruction streams vs # of data streams

Flynn's Classical Taxonomy

- A matrix defining the 4 possible classifications according to Flynn

S I S D Single Instruction stream Single Data stream	S I M D Single Instruction stream Multiple Data stream
M I S D Multiple Instruction stream Single Data stream	M I M D Multiple Instruction stream Multiple Data stream

Single Instruction, Single Data (SISD)


- A serial (non-parallel) computer
- **Single Instruction**
 - Only one instruction stream is being acted on by the CPU during any one clock cycle
- **Single Data**
 - Only one data stream is being used as input during any one clock cycle
- Deterministic execution

Single Instruction, Single Data (SISD)

- This is the **oldest** type of computer
 - **von Neumann system**
 - **Uniprogramming** model
 - Single instruction at a time
 - Fetching and storing one data item at a time
- Examples
 - Older generation mainframes
 - Minicomputers
 - Workstations
 - Single processor/core PCs



Single Instruction, Single Data (SISD)



Single Instruction, Multiple Data (SIMD)


- A type of parallel computer
- **Single Instruction**
 - All processing units execute the same instruction at any given clock cycle
- **Multiple Data**
 - Each processing unit can operate on a different data element
- Same instruction to multiple data items
 - Abstraction
 - System having single control unit and multiple ALUs
- Best suited for specialized problems characterized by a high degree of regularity
 - Graphics/image processing

Single Instruction, Multiple Data (SIMD)

- **Synchronous** (lockstep) and **deterministic** execution
- Two varieties
 - **Processor Arrays** and **Vector Pipelines**
- Examples
 - Processor Arrays
 - Thinking Machines CM-2, MasPar MP-1 & MP-2, ILLIAC IV
 - Vector Pipelines
 - IBM 9000, Cray X-MP, Y-MP & C90, Fujitsu VP, NEC SX-2, Hitachi S820, ETA10

Single Instruction, Multiple Data (SIMD)

- Most modern computers
 - Particularly those with graphics processor units (GPUs) employ SIMD instructions and execution units.



SIMD Example

- Arithmetic operations on two arrays of n elements
 - `for (i = 0; i < n ; i ++)`
 - `x[i] += y[i];`
- SIMD system has n ALUs
 - $x[i]$ and $y[i]$ in ith ALU → storing result in $x[i]$
- m ALUs, where $m < n$
 - Additions in blocks of m elements at a time
 - Ex → $m= 4$ and $n = 15$
- Idling ALUs degrades performance
 - `for (i = 0; i < n ; i ++)`
 - `if (y[i] > 0.0) x[i] += y[i];`
 - Each element of y must be loaded in an ALU for testing before proceeding with the additions

Classical SIMD Systems

- Each ALU operates synchronously
 - Each unit must wait for the next instruction to be broadcast before proceeding
- ALUs have no instruction storage
 - Instruction cannot be stored for later execution
- SIMD systems → ideal for parallelizing simple loops
 - Large data arrays → split among many processors
 - Data-parallelism
 - Very efficient with large data parallel problems
 - Not very efficient with other parallel problems
- In early 1990s → **vector processors**
- Nowadays → GPUs

Vector Processors

- Vectors of data
 - CPUs operating on individual data elements (**scalars**)
- Characteristics
 - **Vector registers**
 - Storing a vector of operands → simultaneous operations
 - Fixed length → 4-128 64-bit elements
 - **Vectorized and pipelined functional units**
 - The same over data → vector operations are **SIMD**
 - **Vector instructions**
 - Instead of scalars → operations on vectors (one single load)
 - **Interleaved memory**
 - Little to no delay in loading/storing successive elements stored over distributed data banks
 - **Strided memory access and hardware scatter/gather**
 - Access in fixed intervals (strides) → Ex: 1st, 5th, 9th ... elements
 - Writing (scatter) or reading (gather) in irregular intervals
 - For both → special hardware to accelerate these operations

Vector Processors




- Virtue
 - Very fast and easy to use for many applications
- Vectorizing compilers
 - Efficient in identifying code for parallelization
 - Determining loops that cannot be parallelized
- Vector systems → high memory bandwidth
 - Every loaded data item is used → not so much caching
 - No handling of irregular data structures
 - Limiting factor for their **scalability** → handling larger problems
 - Limits of operations over longer vectors
 - Commodity systems → very short vectors
 - Custom manufactured → long vectors

Graphics Processing Units

- Real-time graphics application programming interfaces
- Graphics processing pipeline
 - Conversion of internal representation into array of pixels
 - Some stages of this pipeline are programmable
 - Behaviour defined by functions
 - Shader functions → short (a few lines of C code)
 - Implicitly parallel
- GPU → **SIMD parallelism**
 - Large number of ALUs in each GPU core
- **Very high rates of data movement**
 - Large amounts of data (images)
 - To avoid stalls on memory access
- **Heavily relying on hardware multithreading**
 - # of threads relates to registers needed by the shader func.
 - GPU → relatively poor performance on small problems
- **GPU not pure SIMD → multiple cores**
 - So independent instruction streams

SIMD Example Code - Java

- Searching words in file tree




SIMD Example Code - C++

Multiple Instruction, Single Data (MISD)

- A type of parallel computer
- **Multiple Instruction**
 - Each processing unit operates on the data independently via separate instruction streams
- **Single Data**
 - A single data stream is fed into multiple processing units
- Few (if any) actual examples of this class of parallel computer have ever existed
- Some conceivable uses might be
 - Multiple frequency filters operating on a single signal stream

Multiple Instruction, Single Data (MISD)



MISD Example

- Cracking an encrypted code
 - Multiple cryptography algorithms attempting to crack a single coded message
- Running multiple different instances of code over the same data input
 - Testing and verification
 - Experimental analysis → statistical validation
- cracking

Multiple Instruction, Multiple Data (MIMD)

- A type of parallel computer
- **Multiple Instruction**
 - Every processor may be executing a different instruction stream
- **Multiple Data**
 - Every processor may be working with a different data stream
- Execution can be
 - Synchronous or asynchronous
 - Deterministic or non-deterministic


MIMD Systems

- A collection of fully **independent processing units** or cores
 - Each unit with its own control unit and ALU
- MIMD systems are usually **asynchronous**
 - Each processor operating at its own pace
- **There is no global clock**
 - No relation between system times on two different processors
 - Unless the programmer imposes certain synchronization

Multiple Instruction, Multiple Data (MIMD)

- Currently
 - **The most common type of parallel computer**
 - Most modern supercomputers fall into this category
- Examples
 - Most current supercomputers
 - Networked parallel computer clusters and "grids"
 - Multi-processor SMP computers
 - Multi-core PCs
- **Note**
 - Many MIMD architectures also include SIMD execution sub-components

Multiple Instruction, Multiple Data (MIMD)



Class Recap

- Parallel Hardware
 - Flynn's Taxonomy
 - SIMD systems
 - MIMD systems

Next Class

- Parallel Hardware
 - Interconnection networks
 - Cache coherence
 - Shared-memory versus distributed-memory