



Parallel Hardware

Parallel Computing – COSC 3P93

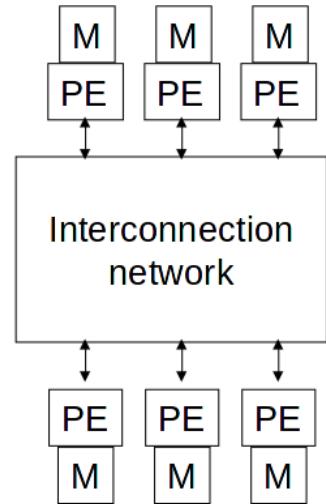
Course Outline

- Introduction
- Parallel Hardware
- Parallel Software
- Performance
- Parallel Programming Models
- Patterns
- SMP with Threads and OpenMP
- Distributed-Memory Programming with MPI
- Algorithms
- Tools
- Parallel Program Design

Today Class

- Parallel Hardware
 - MIMD
 - Distributed Memory
 - Interconnection networks
 - Cache coherence
 - Shared-memory versus distributed-memory

Distributed-memory System



- Multiple computers/nodes
- Message-passing network
- Local memories are private with their own program and data
- **No memory contention**
 - The number of processors is very large
- The processors are connected by
 - **Communication lines**
 - The precise way in which the lines are connected is called the **topology** of the multicomputer
- A typical program consists of subtasks residing in all the memories

Distributed-memory Systems

- **Clusters**
 - Collection of commodity systems → PCs
 - Interconnected by a commodity network → Ethernet
- **Nodes** of a cluster
 - Individual computational units
 - Containing shared-memory systems with multiple multicore processors
 - Joined together by the communication network
- Clusters are not pure distributed-memory systems
 - They may be called **hybrid systems**
 - Nowadays → a cluster will have shared-memory nodes
- **Grid**
 - Infrastructure to support large networks of geographically distributed computers as a unified distributed-memory system
 - They are **heterogeneous**
 - Individual nodes maybe have different types of hardware

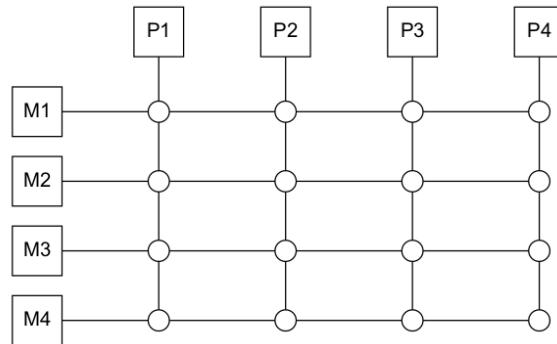
Interconnection Networks

- **Essential** → All communication through them
- Interconnection
 - Great responsible for overall **performance degradation**
- Interconnects have several similarities
 - However → enough particular distinctions to treat them differently
 - **Shared-memory interconnects**
 - **Distributed-memory interconnects**

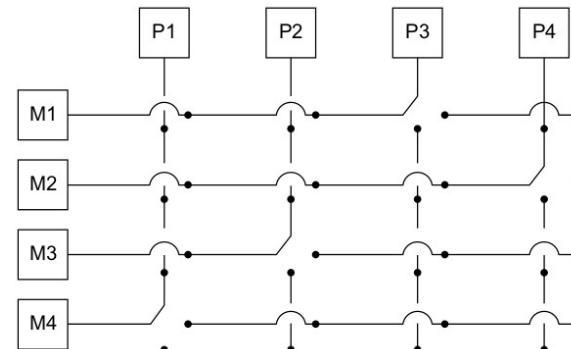
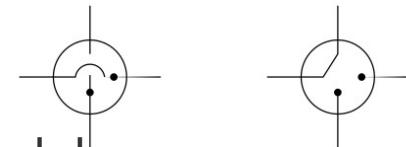
Shared-memory Interconnects

- Widely used interconnects → bus and crossbar
 - **Bus** → parallel communication wires
 - Shared among devices
 - Low cost and flexibility
 - Multiple devices connected at low cost
 - However
 - More connected devices, higher **contention**
 - Lower performance
 - Not good for large-scale shared-memory systems
 - Processors frequently waiting to access the bus

Shared-memory Interco



- Widely used interconnects → bus and crossbar
 - Crossbar
 - Bidirectional comm. links, cores, memory modules, and **switches**
 - **Switched** interconnect → routing of data among devices
 -
 -
 - Simultaneous comm. among different devices
 - Conflicts of two processors accessing the same mem. module
 - Downside → high cost
 - Example
 - P1 writes to M4
 - P2 reads from M3
 - P3 reads M1
 - P4 writes to M2



Distributed-memory Interconnects

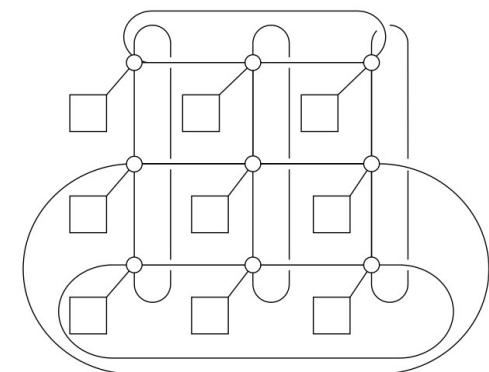
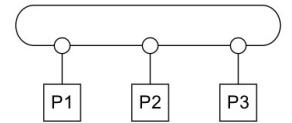
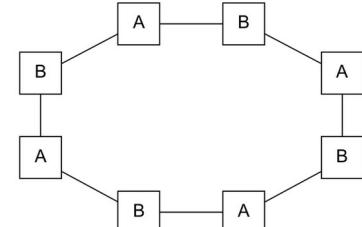
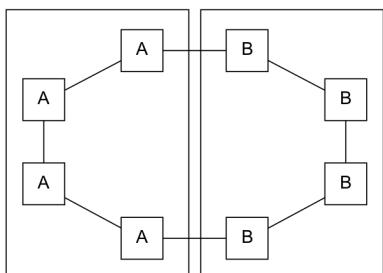
- Directly related to Interconnection Network
- Delimiting factors
 - **Mode of Operation**
 - Synchronous vs. Asynchronous
 - **Control Strategy**
 - Centralized vs. Decentralized
 - **Switching Techniques**
 - Packet switching vs. Circuit switching
 - **Topology**
 - Static vs. Dynamic

Distributed-memory Interconnects

- Two groups → **indirect** and **direct interconnects**

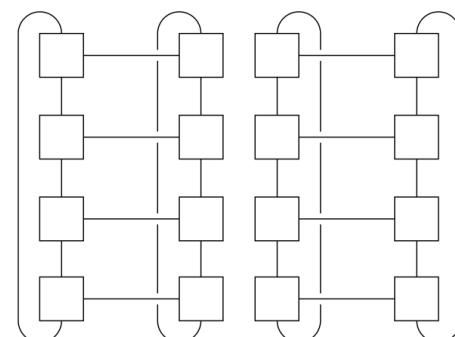
- **Direct interconnect**

- Directly connected to a processor-memory pair
 - Switches connected to each other
 - Examples ($p = \#$ of processors)
 - Ring → $2p$ links
 - 2D toroidal mesh → $3p$ links
 - **Connectivity** factor → # of simultaneous communications
 - **Bisection width**
 - Splitting the system into two equal halves (#nodes)
 - Ex → split the ring topology
 - Two ways of splitting
 - One gives two communications
 - Another gives four communications
 - Incorrect



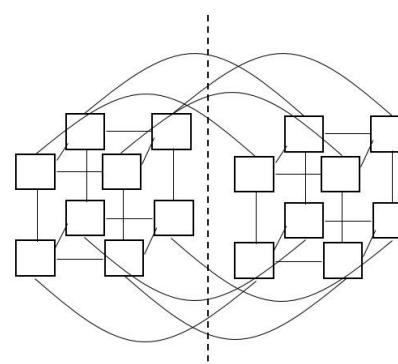
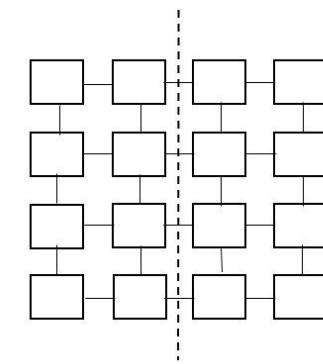
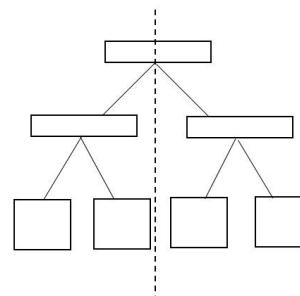
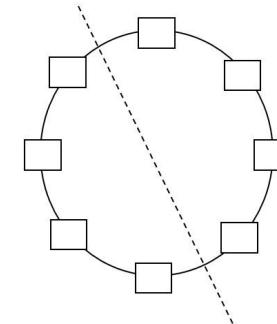
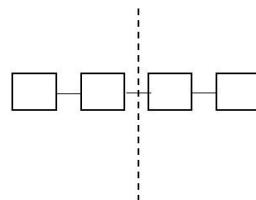
Bisection Width

- It is supposed to give the worst case estimate
 - Previous example → bisection width of ring topology is 2
- **Alternative way of computing bisection width**
 - Setting the minimum number of links needed to be removed to splitting the topology in two equal halves
 - # of links equals to the bisection width
 - Ex → square 2D toroidal mesh → $p = \{q^2 \mid q \text{ is even}\}$
 - Removing the middle horizontal links and respec. Wraparounds
 - Bisection width => $2q = 2\sqrt{p}$



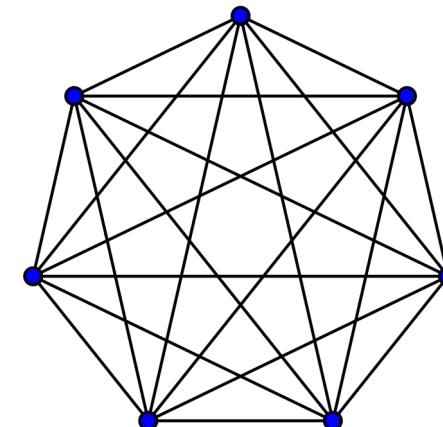
Bisection Width

- Linear array
 - 1 link
- Ring topology
 - 2 links
- Tree topology
 - 1 link
- Mesh topology
 - \sqrt{n} links
- Hyper-cube
 - $n/2$ links



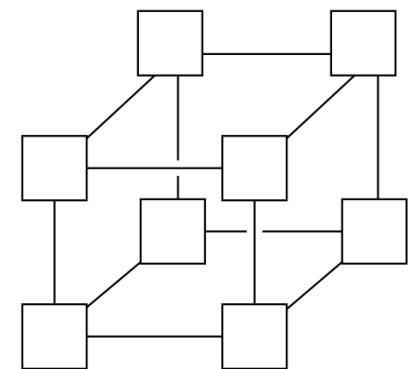
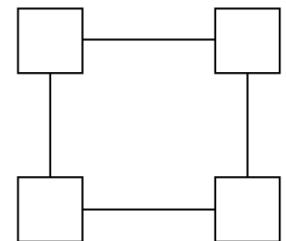
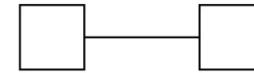
Bandwidth

- Data transmission rate → Mbits or Mbytes per second
- **Bisection bandwidth** → network quality metric
 - Not just the number of links
 - Sum of bandwidths
- Ideal direct interconnect → **fully connected network**
 - Each switch directly connected with every other switch
 - Bisection width → $p^2/4$
 - However
 - Impractical build → $p^2/2 + p/2$
 - But → theoretical best build



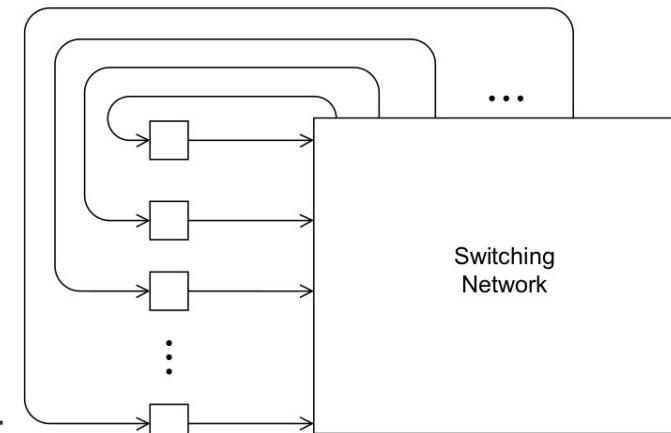
Hypercube

- Direct interconnected
 - Highly connected
- 1D hypercube
 - Fully-connected system
- 2D hypercube
 - 2 1D hypercubes “connected”
- 3D hypercube
 - 2 2D hypercubes “connected”
- Hypercube of dimension d
 - $p = 2^d$
 - Bisection width = $p/2$
 - Each switch supports $1+d = 1 \log_2 p$ wires



Indirect Interconnects

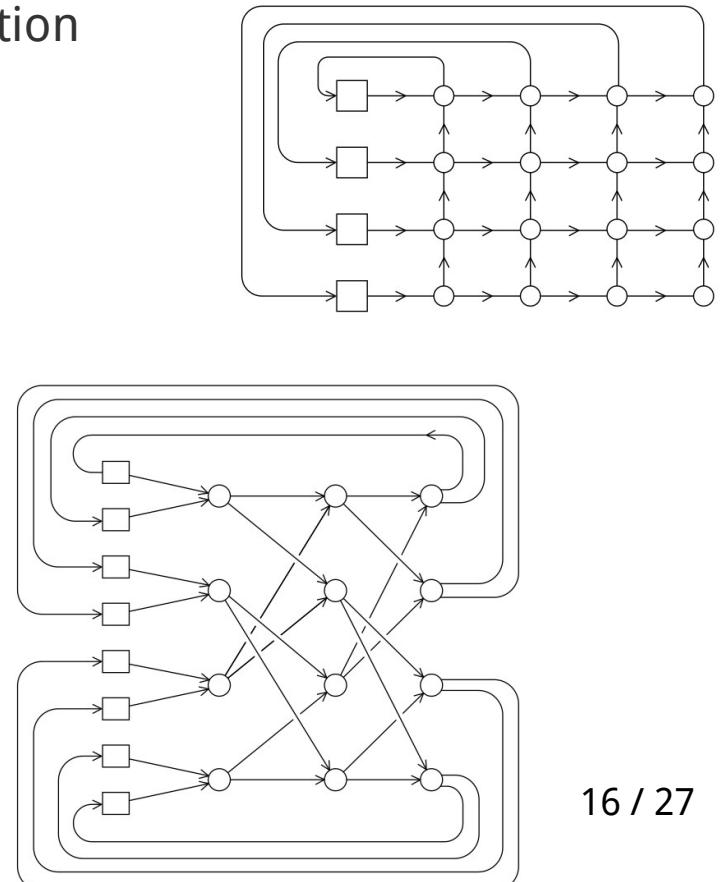
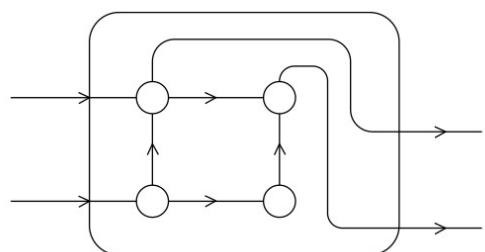
- **Alternative to direct connects**
- Switches not directly connected to processors
 - Each has an incoming, an outgoing link, and a switching network



- More difficult to define bisection width
 - Same principle
 - Dividing the system in two halves
 - The minimum number of links removed so that the two halves do not communicate

Indirect Interconnects Examples

- Crossbar and omega network → for distributed memory
 - **Crossbar** with unidirectional links
 - No 2 procs. attempting to comm. with same proc.
 - Support to simultaneous communication
 - p^2 switches
 - **Omega** network
 - Two-by-two crossbars
 - Less simultaneous communications
 - But, less expensive
 - $2p \log_2 p$ switches



Latency and Bandwidth

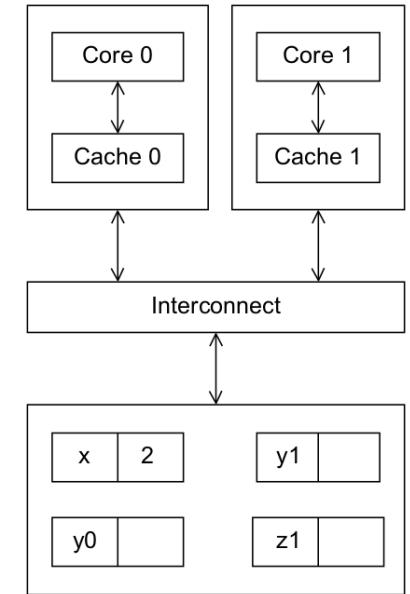
- Time it takes for data to reach its destination
 - Data transmission between
 - Main memory and cache
 - Cache and register
 - HD and memory
 - Between nodes
- **Latency**
 - Elapsing time between beginning of data transmission at source and beginning data receipt at destination → 1st byte
- **Bandwidth**
 - Rate at wich data is received at the destination after the 1st byte is received
- Thus → **message transmission time**
 - $I = n / b$
 - I = latency in seconds
 - b = bandwidth in bytes per sec
 - n = message size in bytes

Latency and Bandwidth

- These terms may be used differently
 - Latency
 - Total message transmission time
 - Or any overhead related to data transmission
 - In other contexts
 - It is not just the time to transmit raw data → it includes
 - Packet headers (encapsulation), CRCs, retransmissions
 - Time to assemble the message and unwrap to the other end
 - Time to disassemble packets and recombine fragmented frames

Cache Coherence

- CPU caches managed by system hardware
 - Important consequences
 - Example
 - Shared-memory system with 2 cores → private data cache each
 - Everything fine while just reading shared data
 - Incoherence happens when modifying shared data → $x = 7$
 - Core 1 may have kept previous value of x in its local cache
 - Unless variable was evicted
 - Not guaranteed that z_1 will be 28 at time 2 in Core 1



Time	Core 0	Core 1
0	$y_0 = x;$	$y_1 = 3*x;$
1	$x = 7;$	Statement(s) not involving x
2	Statement(s) not involving x	$z_1 = 4*x;$

Cache Coherence

- Unpredictability regardless of data caching write-through or write-back policies
 - Write-through policy
 - Main memory will be updated with $x = 7$
 - However, no effect on the local cache of Core 1
 - Write-back policy
 - $X = 7$ in Core 0 will not even make to main memory and be available to Core 1 when z_1 is updated
- Clearly, a major problem → lack of control
 - Apparently correct statements may lead to inconsistency
 - No mechanism for ensuring caching of multiple processors
 - **Cache coherence** problem
 - Two approaches
 - **Snooping cache coherence**
 - **Directory-based cache coherence**

Snooping Cache Coherence

- Approach from bus-based systems
 - Any signal transmitted on the bus is seen by all cores connected to it
- Example
 - Core 0 updating shared variable → broadcasts it
 - Core 1 snoops the bus and catches the updates
- Actually
 - The snooping protocol broadcasts the cache line containing x and not x
- Characteristics
 - Cores do not need to be interconnected by a bus
 - It works fine with write-through and write-back
 - Just additional steps for write-back

Directory-based Cache Coherence

- Problem → Broadcasting is not always possible
 - Large-scale systems
- Use of a **directory**
 - Distributed data structure
 - Storing the status of each cache line
- Characteristics
 - Substantial additional storage needed
 - Updates restrict to contacting only related cores
- Example
 - Core 0 reads a line into its local cache
 - Directory is updated → core 0 has a copy
 - If variable is updated → directory is consulted
 - Cache controllers of the cores containing copies of the line are invalidated

False Sharing

- Reminder → CPU caches implemented in hardware
 - Operating over lines and not variables
 - Consequences on performance
- Example → iterative calls to $f(i,j)$
 - Code can be parallelized by # of cores (core_count)

```
int i , j , m , n ;
double y [ m ];
/* Assign y = 0 */

for ( i = 0; i < m ; i ++)
    for ( j = 0; j < n ; j ++)
        y [ i ] += f ( i , j );
```

```
/* Private variables */
int i, j, iter_count;
/* Shared variables initialized by one core */
int m, n, core_count;
double y [ m ];
iter_count = m / core_count

/* Core 0 does this */
for ( i = 0; i < iter_count; i++)
    for ( j = 0; j < n; j++)
        y [ i ] += f ( i, j );

/* Core 1 does this */
for ( i = iter_count+1; i < 2*iter_count; i++)
    for ( j = 0; j < n; j++)
        y [ i ] += f ( i, j );
. . .
```

False Sharing

- Assume
 - $m = 8$, a large $n \dots$
 - Two cores \rightarrow 0 and 1
 - 8-byte doubles \rightarrow y occupying whole cache line
 - And a directory-based coherence
- What would happen if both cores run simultaneously?
 - Every time they execute $y[i] = f(i, j)$
 - The whole cache line will be invalidated \rightarrow update
 - Next time statement is run, system will have to fetch whole line from memory
 - However \rightarrow core 0 and 1 never access each other's elements of y
- This condition is **false sharing**
 - It does not affect coherence
 - It affects **performance**
- How to solve (partially) the problem?
 - The use of temporary storage

Shared-memory Versus Distributed-memory

- Why so many MIMD systems being shared-memory?
 - Most programmers find the implicit coordination of processors through shared data structure more appealing than explicitly sending messages
- **Major disadvantages of shared-memory**
 - Cost of scaling the interconnect
 - Large crossbars are expensive
 - Dramatic increase of conflicts as scale grows
 - They are suitable for systems with a few processors
 - Distributed-memory interconnects “inexpensive”
 - Hypercube and toroidal mesh
 - Better suited for systems with vast amounts of data and computation
 - Many thousand-processor systems been already built

Class Recap

- Parallel Hardware
 - MIMD
 - Distributed Memory
 - Interconnection networks
 - Cache coherence
 - Shared-memory versus distributed-memory

Next Class

- Parallel Software
 - Caveats
 - Coordinating the processes/threads
 - Shared-memory
 - Distributed-memory
 - Programming hybrid systems