



Parallel Software

Parallel Computing – COSC 3P93

Course Outline

- Introduction
- Parallel Hardware
- **Parallel Software**
- Performance
- Parallel Programming Models
- Patterns
- SMP with Threads and OpenMP
- Distributed-Memory Programming with MPI
- Algorithms
- Tools
- Parallel Program Design

Today Class

- Parallel Software
 - Caveats
 - Coordinating the processes/threads
 - Shared-memory
 - Distributed-memory
 - Programming hybrid systems

Parallel Software

- Just particular class of systems are parallel
 - Operating systems, DBMS, Web servers
- Very little commodity software making extensive use of parallel hardware
- We cannot rely on hardware and compilers for performance
 - Applications need to be designed to exploit shared- and distributed memory architectures
- Terminology
 - Shared-memory systems → forking **threads**
 - Distributed-memory systems → **processes** and **tasks**
 - When in both systems (hybrid) → processes/threads

MIMD and SIMD

- MIMD
 - APIs of GPU grow at rapid pace → they are different from MIMD APIs
- SIMD / SPMD
 - Instead of running different programs on each core
 - Single executable behaving differently according to data
 - Use of conditional branches

```
if ( I'm thread/process 0)
    do this;
else
    do that;
```
- Note that
 - Program is **task parallel** → parallelism is by dividing tasks among threads or processes
- multiprocessing

Coordinating Processes/threads

- High performance → trivial in a few cases
 - Example: addition of 2 arrays

```
double x [ n ], y [ n ];
.
.
for (int i = 0; i < n ; i++)
    x [ i ] += y [ i ];
```
- Parallelizing it
 - Just simple assignment of array elements to processes
 - Process 0 → 0,..., n/(p-1)
 - 1. Splitting the work
 - Same (even) amount of work
 - **Load balancing** → not knowing in advance the workload
 - Converting serial to parallel design → **parallelization**
 - Programs that are turned parallel by simply dividing work → **embarrassingly parallel**
 - Usually parallelization is more complex
 - Requiring steps (2) and (3)
 - 2. Arrange processes/threads to synchronize
 - 3. Arrange communication among them

Shared Memory

- Variables/data can be shared or private
 - **Shared**
 - Read or written by any thread
 - **Private/local**
 - Only one thread accesses it
- Shared variables are essential for communication/interactions among threads
 - Implicit rather than explicit

Dynamic and Static Threads

- **Dynamic threads**
 - Master/slave paradigm
 - Master thread with a collection of empty worker threads
 - Master waiting for work requests
 - Upon a request → master forks a worker thread
 - Upon work termination → worker thread joins the master
 - Efficient use of resources
 - Only being used/reserved while the thread is running

Dynamic and Static Threads

- **Static threads**
 - master/slave paradigm
 - Master thread a with pool of forked (ready to go) threads
 - Threads run while the whole work is finished
 - Threads join master → master may do some cleanup
 - Less efficient use of resources
 - Threads reserve resources while idling
 - It avoids the time-consuming forking and joining operations
 - If resources are available
 - Static paradigm is more efficient → less delays
 - Very close to the mindset of distributed-memory programming
 - Static paradigm is preferred
 - ThreadPoolServer

Nondeterminism

- MIMD → asynchronous processor executions
 - Nondeterminism
 - Input can result in different outputs
 - Multiple independent thread runs:
 - Completing statements at different rates
 - Example → two threads and variable my_x
 - `printf ("Thread%d > my_val = %d\n", my_rank, my_x);`
 - Outputs
 - Thread 0 > my val = 7
 - Thread 1 > my val = 19
 - Or
 - Thread 1 > my val = 19
 - Thread 0 > my val = 7

Nondeterminism

- Nondeterminism is not a problem in many cases
 - Ordering is not a problem or can be corrected
- In many other cases → specially shared-memory
 - It leads to disastrous results → program errors
- Example → threads computing an int
 - Values stored in private var. my_val and shared var. x (0)
 - `my_val = Compute_val (my_rank);`
 - `x += my_val ;`
 - One possible execution → not as expected

Time	Core 0	Core 1
0	Finish assignment to my_val	In call to Compute_val
1	Load x = 0 into register	Finish assignment to my_val
2	Load my_val = 7 into register	Load x = 0 into register
3	Add my_val = 7 to x	Load my_val = 19 into register
4	Store x = 7	Add my_val to x
5	Start other work	Store x = 19

- From attempting to update the memory location x at the same time
- Nondeterminism.BadThreads

Race Condition

- Threads and processes in a “horse race”
 - Outcome depends on the ordering → who wins it
 - In the example → $x += \text{my_val}$
- **Critical section**
 - A block of code that can only be executed by one thread at a time
- **Mutually exclusive** access to the section
 - Program design's responsibility
- Mechanisms → **mutual exclusion lock, mutex, or lock**
 - Support of underlying hardware
 - Critical section protected by a lock
 - Lock should be obtained for granting access to the section

Mutex Example

- Locks could be applied to the race condition example
 - `my_val = Compute_val (my_rank);`
 - `lock (& add my_val_lock);`
 - `x += my_val ;`
 - `unlock (& add_my_val_lock);`
- There is no predetermined order imposed to the threads
 - It just guarantees mutual exclusion
- Mutex downside
 - **Serialization** → one thread accessing at a time
 - **Parallelization contention**
- Thus
 - A code should have as few critical sections as possible
 - Critical sections as short as possible

Busy Waiting

- A loop whose sole purpose is to test a condition
 - Example
 - `my_val = Compute_val (my_rank);`
 - `if (my_rank == 1)`
 - `while (! ok for 1); /* Busy-wait loop */`
 - `x += my_val; /* Critical section */`
 - `if (my_rank == 0)`
 - `ok for 1 = true ; /* Let thread 1 update x */`
 -
- Wasteful method → waste of system resources for testing
 - No useful work is done and core runs thread

Semaphores

- A method for signalling processes and threads
 - Granting access to critical sections
- Different from mutexes
 - Semantically similar
- **Semaphores** allow types o thread synchronization that are easier than using mutexes
- **Monitor**
 - Higher-level mutual exclusion
 - Object-like entity whose methods can only be executed by one thread at a time
- Dining philosophers

Transactions

- Commonly seen in DBMS
- To guarantee consistency and integrity
- All or nothing
 - A transaction is successful
 - If all involved steps are successful
- An error
 - **Rollback** the transaction
- **Transactional memory**
 - In shared-memory programs as transactions
 - Either a thread succeeded completely in a critical section
 - Or any partial results are rolled back
 - Critical section is repeated
- Example → Two-phase locking

Thread Safety

- **Thread safe**
 - Block of code that does not lead to unexpected results when used in multithread program → at any circumstance
 - Threads are nondeterministic by nature
- Example
 - In many cases, parallel programs can call functions developed for use in serial programs
 - Functions with declared local variables are allocated in stack
 - A thread has its own call stack → these are private variables
 - However → C allows declaration of static var. in functions
 - They persist from one call to the next
 - Shared among function callers
 - Practical example → strtok in C library (split string in tokens)
 - Static char* for subsequent calls
 - strtok is **not thread safe** → (strtok_example.cpp)

Distributed Memory

- In distributed memory
 - Cores directly access their own, private memories
- Sharing data → APIs
 - **Message passing**
- Distributed-memory APIs can be used in shared-memory hardware
 - Logically partition memory in private address spaces for threads
 - A library implement needed communication
- Distributed-memory programs composed of processes rather than threads
 - Usually such programs spans over independent CPUs with independent operating systems
- message_passing

Message Passing

- API at minimum → **send** and **receive** functions
- Processes using unique IDs → ranks (0, 1, ..., p-1)
 - Example of message passing

```
char message [100];
.
.
.
my rank = Get rank ();
if ( my rank == 1) {
    sprintf ( message , "Greetings from process 1");
    Send ( message , MSG CHAR , 100, 0);
} else if ( my rank == 0) {
    Receive ( message , MSG CHAR , 100, 1);
    printf ("Process 0 > Received: %s\n", message );
}
```

- Note
 - The program segment is SPMD
 - Two processes following the same executable

Message Passing

- Behaviour of **Send** and **Receive**
 - It depends on the API implementation
 - The simplest
 - A **blocking** *Send* call
 - It blocks until *Receive* starts
 - **Nonblocking** Send
 - Contents of the message placed into storage
 - Commonly → *Receive* **blocks** until caller receives
- Additional functions
 - **Broadcast** → collective communication
 - **Reduction** → combining results of multiple processes
- Widely used API → **Message-Passing Interface (MPI)**
 - Powerful and versatile → used by most powerful computers
 - Very low level → programmer needs to handle a lot of details
 - Called → the assembly language of parallel programming

One-sided Communication

- Usual communication
 - Two acting parts (send and receive)
- **One-sided communication** → remote memory access
 - Single process calling a function
 - Updating or retrieving value from memory
 - Local or remote memory
 - A simplification → just one process
 - Less overhead → less synchronization
 - Less function calls
- Example
 - Process 0 copies value into memory of process 1
 - 0 has ways to know it is safe to copy → overwriting
 - Synchronization of the processes
 - 1 has ways to know about the update
 - Flag variable → process 1 **polling** the flag
 - Busy checking it until flags shows update is complete
- **No direct interactions**
 - Overhead and introduce errors (hard to trace)

Partitioned Global Address Space Languages

- Some may prefer shared-memory programming
 - Instead of message passing or one-sided communication
- Shared memory techniques in programming distributed-memory hardware
 - **Not simple** → example
 - Just treating multiple distributed memory systems as a large memory
→ poor or unpredictable performance
 - Access remote memory may take hundreds or thousands times longer than local memory

```
shared int n = ...;
shared double x[n], y[n];
private int i, my_first_element, my_last_element;
my first element = ...;
my last element = ...;
/* Initialize x and y */

.
.
.

for ( i = my_first_element; i <= my_last_element; i++)
    x[i] += y[i];
```

- All x in core 0 and all y in core 1

Partitioned Global Address Space Languages

- **Partitioned global address space**
 - With tools to avoid delays and performance loss
 - Private variables allocated in local memory of the core running the code
 - Shared data structures used for programmer to control

Programming Hybrid Systems

- Merging/mixing shared-memory and distributed-memory systems
- Allowing the best of both worlds
 - For programs requiring **high performance levels**
- **High complexity** → extremely difficult design
 - Application-specific

Input and Output

- Input and Output calls
 - **Major concern in parallel computing**
- High performance computing
 - **Minimizing I/O calls**
 - Or dealing with **little amount of data**
 - Not the case of Big Data Analytics
- Even the limited use of I/O functions cause problems
 - I/O language functions
 - No details about how they are performed
 - Example → printf and scanf in C, which is a serial language
 - A single process' threads share stdin, stdout, and stderr
 - Simultaneous access to I/O
 - Nondeterministic outcome → not predictable

Class Recap

- Parallel Software
 - Caveats
 - Coordinating the processes/threads
 - Shared-memory
 - Distributed-memory
 - Programming hybrid systems

Next Class

- Performance
 - Speed up