

Performance

Parallel Computing – COSC 3P93

Course Outline

- Introduction
- Parallel Hardware
- Parallel Software
- **Performance**
- Parallel Programming Models
- Patterns
- SMP with Threads and OpenMP
- Distributed-Memory Programming with MPI
- Algorithms
- Tools
- Parallel Program Design


Today Class

- Performance
 - Performance scalability
 - Analytical performance measures
 - Amdahl's law and Gustafson-Barsis' law

What is Performance?

- In computing, performance is defined by 2 factors
 - Computational requirements (what needs to be done)
 - Computing resources (what it costs to do it)
- Computational problems translate to requirements
- Computing resources interplay and tradeoff

$$\text{Performance} \sim \frac{1}{\text{Resources for solution}}$$



Hardware



Time



Energy

... and ultimately



Money

Why do we care about Performance?

- Performance itself is a measure of how well the computational requirements can be satisfied
- We evaluate performance to understand the relationships between requirements and resources
 - Decide how to change “solutions” to target objectives
- Performance measures reflect decisions about how and how well “solutions” are able to satisfy the computational requirements

“The most constant difficulty in contriving the engine has arisen from the desire to reduce the time in which the calculations were executed to the shortest which is possible.”

Charles Babbage, 1791 – 1871

What is Parallel Performance?

- Here we are concerned with performance issues when using a parallel computing environment
 - Performance with respect to parallel computation
- Performance is the *raison d'être* for parallelism
 - Parallel performance versus sequential performance
 - If the **performance** is not better
 - Parallelism is not necessary
- Parallel processing includes techniques and technologies necessary to compute in parallel
 - Hardware, networks, operating systems, parallel libraries, languages, compilers, algorithms, tools, ...
- Parallelism must deliver performance
 - How? How well?

Performance Expectation (Loss)

- If each processor is rated at k MFLOPS and there are p processors, should we see $k \cdot p$ MFLOPS performance?
- If it takes 100 seconds on 1 processor
 - Shouldn't it take 10 seconds on 10 processors?
- Several causes affect performance
 - Each must be understood separately
 - But they interact with each other in complex ways
 - Solution to one problem may create another
 - One problem may mask another
- Scaling (system, problem size) can change conditions
- Need to understand performance space

Embarrassingly Parallel Computations

- An embarrassingly parallel computation is one that can be obviously divided into completely independent parts that can be executed simultaneously
 - In a truly embarrassingly parallel computation there is no interaction between separate processes
 - In a nearly embarrassingly parallel computation results must be distributed and collected/combined in some way
- Embarrassingly parallel computations have potential to achieve maximal speedup on parallel platforms
 - If it takes T time sequentially, there is the potential to achieve T/P time running in parallel with P processors
 - What would cause this not to be the case always?

Scalability

- A program can scale up to use many processors
 - What does that mean?
- How do you evaluate scalability?
- How do you evaluate scalability goodness?
- Comparative evaluation
 - If double the number of processors, what to expect?
 - Is scalability linear?
- Use parallel efficiency measure
 - Is efficiency retained as problem size increases?
- Apply performance metrics

Performance and Scalability

- Evaluation
 - **Sequential** runtime (T_{seq}) is a function of
 - Problem size and architecture
 - **Parallel** runtime (T_{par}) is a function of
 - Problem size and parallel architecture
 - # processors used in the execution
 - Parallel performance affected by
 - Algorithm + architecture
- **Scalability**
 - Ability of parallel algorithm to achieve performance gains proportional to the number of processors and the size of the problem

Scalability

- In general → a problem is **scalable**
 - If it can handle ever increasing problem sizes
- The problem is **strongly scalable**
 - If we increase the number of processes/threads and keep the efficiency fixed without increasing problem size
- The problem is **weakly scalable**
 - If we keep the efficiency fixed by increasing the problem size at the same rate as we increase the number of processes/threads

Taking Timings



- What is time?
- Start to finish?
- A program segment of interest?
- CPU time?
- Wall clock time?

Performance Metrics and Formulas

- T_1 is the execution time on a single processor
- T_p is the execution time on a N-processor system
- $S(p)$ (S_p) is the speedup
- $E(p)$ (E_p) is the efficiency
- Cost(p) (C_p) is the cost

$$S(p) = \frac{T_1}{T_p}$$

$$\text{Efficiency} = \frac{S_p}{N} = \frac{\left(\frac{T_{seq}}{T_{par}}\right)}{N} = \frac{T_{seq}}{N \cdot T_{par}}$$

$$\text{Cost} = N \times T_p$$

- Parallel algorithm is cost-optimal
 - Parallel time = sequential time ($C_p = T_1$, $E_p = 100\%$)



Speedups and Efficiency

- Parallel program

p	1	2	4	8	16
S	1.0	1.9	3.6	6.5	10.8
$E = S/p$	1.0	0.95	0.90	0.81	0.68

- Parallel program on different problem sizes

	p	1	2	4	8	16
Half	S	1.0	1.9	3.1	4.8	6.2
	E	1.0	0.95	0.78	0.60	0.39
Original	S	1.0	1.9	3.6	6.5	10.8
	E	1.0	0.95	0.90	0.81	0.68
Double	S	1.0	1.9	3.9	7.5	14.2
	E	1.0	0.95	0.98	0.94	0.89



Speedups and Efficiency

- Parallel program on different problem sizes

	p	1	2	4	8	16
Half	S	1.0	1.9	3.1	4.8	6.2
	E	1.0	0.95	0.78	0.60	0.39
Original	S	1.0	1.9	3.6	6.5	10.8
	E	1.0	0.95	0.90	0.81	0.68
Double	S	1.0	1.9	3.9	7.5	14.2
	E	1.0	0.95	0.98	0.94	0.89

Limits and Costs of Parallel Programming

- **Amdahl's Law**
 - Amdahl's Law states that potential program speedup is defined by the fraction of code (P) that can be parallelized
$$speedup = \frac{1}{1-P}$$
 - If none of the code can be parallelized
 - P = 0 and the speedup = 1 (no speedup)
 - If all of the code is parallelized
 - P = 1 and the speedup is infinite (in theory)
 - If 50% of the code can be parallelized
 - Maximum speedup = 2
 - The code will run twice as fast
 - Introducing the number of processors performing the parallel fraction of work
 - The relationship can be modelled by where P = parallel fraction, N = number of processors and S = serial fraction

Amdahl's Law (Fixed Size Speedup)


- Let S be the fraction of a program that is sequential
 - $1-S$ is the fraction that can be parallelized
- Let T_1 be the execution time on 1 processor
- Let T_p be the execution time on N processors
- S_p is the speedup
 - $S_p = T_1 / T_p$
 - $= T_1 / (S \cdot T_1 + (1-S) \cdot T_1 / N)$
 - $= 1 / (S + (1-S) / N)$
- As $N \rightarrow \infty$
 - $S_p = 1 / S$

Limits and Costs of Parallel Programming

- Amdahl's Law
 - Simplifying

$$Sp = \frac{1}{\frac{P}{N} + S}$$

- where P = parallel fraction
- N = number of processors
- S = serial fraction



N	Speedup			
	P=0.50	P=0.90	P=0.95	P=0.99
10	1.82	5.26	6.89	9.17
100	1.98	9.17	16.80	50.25
1000	1.99	9.91	19.62	90.99
10000	1.99	9.91	19.96	99.02
100000	1.99	9.99	19.99	99.90

Limits and Costs of Parallel Programming

- **Amdahl's Law**



- It soon becomes obvious that there are limits to the scalability of parallelism
 - For example

N	Speedup			
	P=0.50	P=0.90	P=0.95	P=0.99
10	1.82	5.26	6.89	9.17
100	1.98	9.17	16.80	50.25
1000	1.99	9.91	19.62	90.99
10000	1.99	9.91	19.96	99.02
100000	1.99	9.99	19.99	99.90

Limits and Costs of Parallel Programming

- **Amdahl's Law**

- It soon becomes obvious that there are limits to the scalability of parallelism
 - For example



Amdahl's Law and Scalability

- Scalability
 - Ability of parallel algorithm to achieve performance gains proportional to the number of processors and the size of the problem
- When does Amdahl's Law apply?
 - **When the problem size is fixed**
 - Strong scaling ($N \rightarrow \infty$, $S_p \rightarrow S_\infty = 1 / S$)
 - Speedup bound is determined by the degree of sequential execution time in the computation, not # processors!!!
 - Uhh, this is not good ... Why?
 - Perfect efficiency is hard to achieve
- ** Effect of overhead!
 - Assume a perfect, virtual “parallelization” of the serial code

$$T_{\text{par}} = T_{\text{seq}} / p + T_{\text{overhead}}$$

Gustafson-Barsis' Law (Scaled Speedup)


- Often interested in larger problems when scaling
 - How big of a problem can be run (HPC Linpack)
 - Constrain problem size by parallel time
- Assume parallel time is kept constant
 - T_p is constant
 - P_{seq} is the fraction of T_p spent in sequential execution
 - P_{par} is the fraction of T_p spent in parallel execution
- What is the execution time on one processor?
 - $T_s = P_{seq} \cdot T_p + (1 - P_{seq}) \cdot N \cdot T_p$
- What is the speedup in this case?
 - $S_p = T_s / T_p$
 - $= (P_{seq} \cdot T_p + (1 - P_{seq}) \cdot N \cdot T_p) / T_p$
 - $= P_{seq} + (1 - P_{seq}) \cdot N$
 - $= (1 - P_{par}) + P_{par} \cdot N$

Gustafson-Barsis' Law and Scalability

- Scalability
 - Ability of parallel algorithm to achieve performance gains proportional to the number of processors and the size of the problem
- When does Gustafson's Law apply?
 - When the problem size can increase as the number of processors increases
 - **Weak scaling** → ($Sp = P_{seq} + NP_{par}$)
 - Speedup function includes the number of processors!!!
 - Can maintain or increase parallel efficiency as the problem scales


Amdahl versus Gustafson-Baris

Amdahl




Amdahl versus Gustafson-Baris

Gustafson-Baris




DAG Model of Computation

- Think of a program as a directed acyclic graph (DAG) of tasks
 - A task can not execute until all the inputs to the tasks are available
 - These come from outputs of earlier executing tasks
 - DAG shows explicitly the task dependencies
- Think of the hardware as consisting of workers (processors)
- Consider a greedy scheduler of the DAG tasks to workers
 - No worker is idle while there are tasks still to execute




Work-Span Model

- T_p = time to run with P workers
- T_1 = work
 - Time for serial execution
 - Execution of all tasks by 1 worker
- Sum of all work
- T_∞ = span
 - Time along the critical path
- Critical path
 - Sequence of task execution (path) through DAG that takes the longest time to execute
 - Assumes an infinite # workers available




Work-Span Example

- Let each task take 1 unit of time
- DAG at the right has 7 tasks
- $T_1 = 7$
 - All tasks have to be executed
 - Tasks are executed in a serial order
 - Can the execution be in any order?
- $T_\infty = 5$
 - Time along the critical path
 - In this case, it is the longest pathlength of any task order that maintains necessary dependencies




Lower/Upper Bound on Greedy Scheduling

- Suppose we only have P workers
- We can write a work-span formula to derive a lower bound on T_P
 - $\text{Max}(T_1 / P, T_\infty) \leq T_P$
- T_∞ is the best possible execution time
- Brent's Lemma derives an upper bound
 - Capture the additional cost executing the other tasks not on the critical path
 - Assume we can do so without overhead
 - $T_P \leq (T_1 - T_\infty) / P + T_\infty$





Consider Brent's Lemma for 2 Processors

- $T_1 = 7$
- $T_\infty = 5$
- $T_2 \leq (T_1 - T_\infty) / P + T_\infty$
 $\leq (7 - 5) / 2 + 5$
 ≤ 6



Amdahl was an optimist!



Estimating Running Time

- Scalability requires that T_∞ be dominated by T_1

$$T_P \approx T_1 / P + T_\infty \text{ if } T_\infty \ll T_1$$

- Increasing work hurts parallel execution proportionately
- The span impacts scalability
 - Even for finite P

Parallel Slack

- Sufficient parallelism implies linear speedup

$$T_P \approx T_1/P \quad \text{if} \quad T_1/T_\infty \gg P$$



Linear speedup



Parallel slack

Asymptotic Complexity

- Time complexity of an algorithm summarizes how the execution time grows with input size
- Space complexity summarizes how memory requirements grow with input size
- Standard work-span model considers only computation, not communication or memory
- Asymptotic complexity is a strong indicator of performance on large-enough problem sizes and reveals an algorithm's fundamental limits

Definitions for Asymptotic Notation

- Let $T(N)$ mean the execution time of an algorithm
- Big O notation
 - $T(N)$ is a member of $O(f(N))$ means that
 - $T(N) \leq c \cdot f(N)$ for constant c
- Big Omega notation
 - $T(N)$ is a member of $\Omega(f(N))$ means that
 - $T(N) \geq c \cdot f(N)$ for constant c
- Big Theta notation
 - $T(N)$ is a member of $\Theta(f(N))$ means that
 - $c_1 \cdot f(N) \leq T(N) < c_2 \cdot f(N)$ for constants c_1 and c_2

Limits and Costs of Parallel Programming

- **Quote**
 - You can spend a lifetime getting 95% of your code to be parallel, and never achieve better than 20x speedup no matter how many processors you throw at it!
- Certain problems demonstrate increased performance by increasing the problem size
 - 2D Grid Calculations 85 seconds 85%
 - Serial fraction 15 seconds 15%
 - We can increase the problem size by doubling the grid dimensions and halving the time step
 - This results in four times the number of grid points and twice the number of time steps
 - 2D Grid Calculations 680 seconds 97.84%
 - Serial fraction 15 seconds 2.16%
 - Problems that increase the percentage of parallel time with their size are more scalable than problems with a fixed percentage of parallel time

Limits and Costs of Parallel Programming

- **Complexity**
 - In general, parallel applications are much more complex than corresponding serial applications, perhaps an order of magnitude
 - Not only there are multiple instruction streams executing at the same time, but you also there are data flowing between them
 - The costs of complexity are measured in programmer time in virtually every aspect of the software development cycle
 - Design
 - Coding
 - Debugging
 - Tuning
 - Maintenance
 - Adhering to **good** software development practices is essential when working with parallel applications
 - Especially if somebody besides you will have to work with the software

Limits and Costs of Parallel Programming

- **Portability**
 - Thanks to standardization in several APIs
 - MPI, POSIX threads, and OpenMP
 - Portability issues with parallel programs are not as serious as in years past → However...
 - The usual portability issues associated with serial programs apply to parallel programs
 - If you use vendor "enhancements" to Fortran, C, or C++
 - Portability will be a problem.
 - Standards exist for several APIs
 - However → implementations will differ in a number of details
 - Requiring code modifications in order to effect portability
 - Operating systems
 - They can play a key role in code portability issues
 - Hardware architectures
 - Characteristically highly variable and can affect portability

Limits and Costs of Parallel Programming

- **Resource Requirements**
 - The primary intent of parallel programming
 - To decrease execution wall clock time
 - However → more CPU time is required
 - For example → a parallel code that runs in 1 hour on 8 processors actually uses 8 hours of CPU time
 - The amount of memory required can be greater for parallel codes than serial codes
 - The need to replicate data and for overheads associated with parallel support libraries and subsystems
 - Short running parallel programs
 - A decrease in performance compared to a similar serial implementation.
 - Overhead costs associated with setting up the parallel environment, task creation, communications, and task termination
 - A significant portion of the total execution time for short runs

Limits and Costs of Parallel Programming


- **Scalability**
 - Two types of scaling based on time to solution: strong scaling and weak scaling
 - **Strong scaling**
 - The total problem size stays fixed as more processors are added
 - Goal is to run the same problem size faster
 - Perfect scaling means problem is solved in $1/P$ time (compared to serial)
 - **Weak scaling**
 - The problem size per processor stays fixed as more processors are added
 - The total problem size is proportional to the number of processors used.
 - Goal is to run larger problem in same amount of time
 - Perfect scaling means problem Px runs in same time as single processor run

Limits and Costs of Parallel Programming

- **Scalability**
 - Ability of a parallel program's performance to scale
 - Result of a number of interrelated factors
 - Simply adding more processors is rarely the answer
 - The algorithm may have inherent limits to scalability
 - Adding more resources causes performance to decrease
 - A common situation with many parallel applications
 - Hardware factors play a significant role in scalability
 - Memory-cpu bus bandwidth on an SMP machine
 - Communications network bandwidth
 - Amount of memory available on any given machine or set of machines
 - Processor clock speed
 - Parallel support libraries and subsystems software can limit scalability independent of your application

Limits and Costs of Parallel Programming

- Scalability



Class Recap

- Performance
 - Performance scalability
 - Analytical performance measures
 - Amdahl's law and Gustafson-Barsis' law

Next Class

- Performance
 - Scalable parallel execution
 - Parallel execution models
 - Isoefficiency