

Performance

Parallel Computing – COSC 3P93

Course Outline

- Introduction
- Parallel Hardware
- Parallel Software
- **Performance**
- Parallel Programming Models
- Patterns
- SMP with Threads and OpenMP
- Distributed-Memory Programming with MPI
- Algorithms
- Tools
- Parallel Program Design

Today Class

- Performance
 - Scalable parallel execution
 - Parallel execution models
 - Isoefficiency

Scalable Parallel Computing

- Scalability in parallel architecture
 - Processor numbers
 - Memory architecture
 - Interconnection network
 - Avoid critical architecture bottlenecks
- Scalability in computational problem
 - Problem size
 - Computational algorithms
 - Computation to memory access ratio
 - Computation to communication ratio
- Parallel programming models and tools
- Performance scalability

Why Aren't Parallel Applications Scalable?

- **Sequential performance**
- **Critical Paths**
 - Dependencies between computations spread across processors
- **Bottlenecks**
 - One processor holds things up
- **Algorithmic overhead**
 - Some things just take more effort to do in parallel
- **Communication overhead**
 - Spending increasing proportion of time on communication
- **Load Imbalance**
 - Makes all processor wait for the “slowest” one
 - Dynamic behaviour
- **Speculative loss**
 - Do A and B in parallel, but B is ultimately not needed

Critical Paths

- **Long chain of dependence**
 - Main limitation on performance
 - Resistance to performance improvement
- **Diagnostic**
 - Performance stagnates to a (relatively) fixed value
 - Critical path analysis
- **Solution**
 - Eliminate long chains if possible
 - Shorten chains by removing work from critical path

Bottlenecks


- **How to detect?**
 - One processor A is busy while others wait
 - Data dependency on the result produced by A
- **Typical situations**
 - N-to-1 reduction / computation / 1-to-N broadcast
 - One processor assigning job in response to requests
- **Solution techniques**
 - More efficient communication
 - Hierarchical schemes for master slave
- **Program may not show ill effects for a long time**
- Shows up when scaling

Algorithmic Overhead

- Different sequential algorithms to solve the same problem
- All parallel algorithms are sequential
 - When run on 1 processor
- All parallel algorithms introduce addition operations (Why?)
 - Parallel overhead
- Where should the starting point be for a parallel algorithm?
 - Best sequential algorithm might not parallelize at all
 - Or, it doesn't parallelize well (e.g., not scalable)
- What to do?
 - Choose algorithmic variants that minimize overhead
 - Use two level algorithms
- Performance is **the rub**
 - Are you achieving better parallel performance?
 - Must compare with the best sequential algorithm


What is the maximum parallelism possible?

- Depends on application, algorithm, program
 - Data dependencies in execution
- Remember MaxPar
 - Analyzes the earliest possible “time” any data can be computed
 - Assumes a simple model for time it takes to execute instruction or go to memory
 - Result is the maximum parallelism available
- Parallelism varies!



Embarrassingly Parallel Computations


- No or very little communication between processes
- Each process can do its tasks without any interaction with other processes



- Examples
 - Numerical integration
 - Mandelbrot set
 - Monte Carlo methods

Calculating Pi with Monte Carlo

- Consider a circle of unit radius
- Place circle inside a square box with side of 2 cm




- The ratio of the circle area to the square area is

$$\frac{\pi * 1 * 1}{2 * 2} = \frac{\pi}{4}$$

Monte Carlo Calculation of Pi

- Randomly choose a number of points in the square
- For each point p
 - Determine if p is inside the circle
- The ratio of points in the circle to points in the square will give an approximation of $\pi/4$



Performance Metrics Review

- T_1 is the execution time on a single processor
- T_p is the execution time on a p processor system
- $S(p)$ (S_p) is the speedup

$$S(p) = \frac{T_1}{T_p}$$

- $E(p)$ (E_p) is the efficiency

$$\text{Efficiency} = \frac{S_p}{p}$$

- Cost(p) (C_p) is the cost

$$\text{Cost} = p \times T_p$$


- Parallel algorithm is cost-optimal
 - Parallel time = sequential time ($C_p = T_1$, $E_p = 100\%$)

Analytical / Theoretical Techniques

- Involves simple algebraic formulas and ratios
 - Typical variables
 - Data size (N), number of processors (P), machine constants
 - To model performance of individual operations, components, algorithms in terms of the above
 - Be careful to characterize variations across processors
 - Model them with max operators
 - **Constants are important in practice**
 - Use asymptotic analysis carefully
- Scalability analysis
 - Isoefficiency (Kumar)

Isoefficiency

- Goal is to quantify scalability
- How much increase in problem size is needed to retain the same efficiency on a larger machine?
- Efficiency
 - $T_1 / (p * T_p)$
 - $T_p = \text{computation} + \text{communication} + \text{idle}$
- Isoefficiency
 - Equation for equal-efficiency curves
 - If no solution
 - Problem is not scalable in the sense defined by isoefficiency



Scalability of Adding n Numbers



- Scalability of a parallel system is a measure of its capacity to increase speedup with more processors
- Adding **n** numbers on **p** processors with strip partition

$$T_{par} = \frac{n}{p} - 1 + 2 \log p$$

$$\text{Speedup} = \frac{\frac{n-1}{n}}{\frac{n-1}{p} + 2 \log p}$$

$$\approx \frac{n}{\frac{n}{p} + 2 \log p}$$

$$\text{Efficiency} = \frac{S}{p} = \frac{n}{n + 2p \log p}$$



Problem Size and Overhead

- Informally, problem size is expressed as a parameter of the input size
- A consistent definition of the size of the problem is the total number of basic operations (T_{seq})
 - Also refer to problem size as work ($W = T_{\text{seq}}$)
- Overhead of a parallel system is defined as the part of the cost not in the best serial algorithm
- Denoted by T_o , it is a function of W and p

$$T_o(W,p) = pT_{\text{par}} - W \quad (pT_{\text{par}} \text{ includes overhead})$$

$$T_o(W,p) + W = pT_{\text{par}}$$

Isoefficiency Function

- With a fixed efficiency, W is as a function of p

$$T_{par} = \frac{W + T_O(W, p)}{p} \quad W = T_{seq}$$

$$\text{Speedup} = \frac{W}{T_{par}} = \frac{Wp}{W + T_O(W, p)}$$


$$\text{Efficiency} = \frac{S}{p} = \frac{W}{W + T_O(W, p)} = \frac{1}{1 + \frac{T_O(W, p)}{W}}$$

$$E = \frac{1}{1 + \frac{T_O(W, p)}{W}} \rightarrow \frac{T_O(W, p)}{W} = \frac{1 - E}{E}$$

$$W = \frac{E}{1 - E} T_O(W, p) = K T_O(W, p) \quad \text{Isoefficiency Function}$$

Isoefficiency Function of Adding n Numbers

- Overhead function:
 - $T_o(W,p) = pT_{par} - W = 2p\log(p)$
- Isoefficiency function
 - $W=K*2p\log(p)$
- If p doubles
 - W needs also to be doubled to roughly maintain the same efficiency
- Isoefficiency functions can be more difficult to express for more complex algorithms



More Complex Isoefficiency Functions

- A typical overhead function T_o can have several distinct terms of different orders of magnitude with respect to both p and W
- We can balance W against each term of T_o and compute the respective isoefficiency functions for individual terms
 - Keep only the term that requires the highest grow rate with respect to p
 - This is the asymptotic isoefficiency function

Isoefficiency

- Consider a parallel system with an overhead function

$$T_O = p^{3/2} + p^{3/4}W^{3/4}$$

- Using only the first term

$$W = Kp^{3/2}$$

- Using only the second term

$$W = Kp^{3/4}W^{3/4}$$

$$W^{1/4} = Kp^{3/4}$$

$$W = K^4 p^3$$

- $K^4 p^3$ is the overall asymptotic isoefficiency function

Parallel Computation (Machine) Models

- **PRAM** (parallel RAM)
 - Basic parallel machine
- **BSP** (Bulk Synchronous Parallel)
 - Isolates regions of computation from communication
- **LogP**
 - Used for studying distributed memory systems
 - Focuses on the interconnection network
- **Roofline**
 - Based in analyzing “feeds” and “speeds”

PRAM Complexity Measures

- For each individual processor
 - **Time**: number of instructions executed
 - **Space**: number of memory cells accessed
- PRAM machine
 - **Time**: time taken by the longest running processor
 - **Hardware**: maximum number of active processors
- Technical issues
 - How processors are activated
 - How shared memory is accessed

PRAM is a Theoretical (Unfeasible) Model

- Interconnection network between processors and memory would require a very large amount of area
- The message-routing on the interconnection network would require time proportional to network size
- Algorithm's designers can forget the communication problems and focus their attention on the parallel computation only
- There exist algorithms simulating any PRAM algorithm on bounded degree networks
- Design general algorithms for the PRAM model and simulate them on a feasible network

Classification of PRAM Models

- EREW (Exclusive Read Exclusive Write)
 - No concurrent read/writes to the same memory location
- CREW (Concurrent Read Exclusive Write)
 - Multiple processors may read from the same global memory location in the same instruction step
- ERCW (Exclusive Read Concurrent Write)
 - Concurrent writes allowed
- CRCW (Concurrent Read Concurrent Write)
 - Concurrent reads and writes allowed
- CRCW > (ERCW,CREW) > EREW

Complexity of PRAM Algorithms


Problem Model	EREW	CRCW
Search	$O(\log n)$	$O(1)$
List Ranking	$O(\log n)$	$O(\log n)$
Prefix	$O(\log n)$	$O(\log n)$
Tree Ranking	$O(\log n)$	$O(\log n)$
Finding Minimum	$O(\log n)$	$O(1)$

BSP Overview

- **Bulk Synchronous Parallelism**
- A parallel programming model
- Invented by Leslie Valiant at Harvard
- Enables performance prediction
- SPMD (Single Program Multiple Data) style
- Supports both direct memory access and message passing semantics
- BSPlib is a BSP library implemented at Oxford

BSP Supersteps

- A BSP computation consists of a sequence of supersteps
- In each superstep
 - Processes execute computations using locally available data
 - And issue communication requests
- Processes synchronized at the end of the superstep
 - At which all communications issued have been completed




BSP Performance Model

Parameters

- p = number of processors
- l = barrier latency, cost of achieving barrier synchronization
- g = communication cost per word
- s = processor speed
- l , g , and s are measured in FLOPS
- Any processor sends and receives at most h messages in a single superstep (called h -relation communication)
- Time for a superstep = max number of local operations performed by any one processor + $g*h + l$

LogP

- Latency in sending a (small) message between modules
- overhead felt by the processor on sending or receiving message
- gap between successive sends or receives ($1/\text{BW}$)
- Processors



LogP "Philosophy"

- Strategy
 - Mapping of N words onto P processors
 - Computation within a processor
 - Its cost and balance
 - Communication between processors
 - Its cost and balance
- Characterize processor and network performance
- Do not think about what happens in the network
- This should be enough

Class Recap

- Performance
 - Scalable parallel execution
 - Isoefficiency
 - Parallel execution models

Next Class

- Parallel Programming Models